

NISHANT SANWARIA

Agentic AI Developer | LangChain / LangGraph / RAG / MCP / Multi-Agent Systems

104 Vijay Sarita State, Panchsheel, Ajmer, Rajasthan | 77373-96662 | sanwarianishant@gmail.com | github.com/nishant264

PROFILE

B.Tech student (AI & Data Science) building autonomous, tool-using AI agents - from deterministically-gated recovery agents and multi-agent interview coaching systems, to corrective RAG pipelines with MCP-based tool-calling. Comfortable across the full agent stack: LLM orchestration, multi-agent design, RAG, vector search, MCP, and structured tool-calling. Seeking an internship to apply and grow this skill set in a production setting.

EDUCATION

B.Tech in Artificial Intelligence & Data Science

Jaipur, Rajasthan

JECRC Foundation

Sep. 2023 – May 2027

Relevant Coursework: Data Structures & Algorithms, Artificial Intelligence, Internet Technology, RAG Pipelines, Agentic AI, LangGraph

PROJECTS

AI Revenue Recovery Agent - Razorpay Buildathon | Python, Groq API, Deterministic Policy Gating

2026

- Designed an autonomous agent pipeline (Detect, Diagnose, Decide, Gate, Execute, Verify, Report) that diagnoses payment-failure root causes and proposes bounded recovery actions via the Groq LLM API.
- Enforced a deterministic policy gate that unconditionally overrides the LLM's decision on fraud-flagged transactions - 100% fraud hold-back rate with 0% agent failsafe errors.
- Built a synthetic 100-scenario evaluation harness with a full per-case audit trail, achieving 71.1% revenue recovery and 78% diagnosis accuracy against ground truth; covered by 39 unit tests.

Code: github.com/nishant264/ai-revenue-recovery-agent

AI Mock Interview Coach | Python, Streamlit, OpenAI/Groq API, Pydantic

2026

- Built a 4-agent adaptive interview system (Interviewer, Evaluator, Planner, Coach) with a deterministic orchestrator that enforces question bounds and follow-up rules independent of LLM drift.
- Implemented Pydantic-validated structured outputs with one-retry repair and deep-copied session state, keeping live interviews resilient to malformed model responses and provider failures.
- Shipped a live Streamlit app with per-dimension scoring charts and a 7-day coaching report, validated by a 24-test deterministic suite driven by a scripted fake LLM client.

Code: github.com/nishant264/ai-mock-interview-coach

Corrective RAG Agent | Python, LangChain, LangGraph, Qdrant, Claude API, Tavily

2026

- Implemented a corrective, multi-stage RAG pipeline in LangGraph combining Qdrant vector retrieval, LLM-based document relevance grading, query rewriting, and a Tavily web-search fallback for low-relevance results.
- Used Claude 4.5 Sonnet for relevance grading and OpenAI embeddings for retrieval, with a Streamlit UI for document upload and querying.

Code: [github.com/nishant264/Corrective- RAG](https://github.com/nishant264/Corrective-RAG)

TECHNICAL SKILLS

AI & Agents: LangChain, LangGraph, Multi-Agent Orchestration (Agency Swarm, Agno), RAG, Corrective RAG, MCP (Model Context Protocol), Deterministic Guardrails / Policy Gating, ChromaDB, FAISS, Qdrant

LLM Tools: Gemini API, OpenAI API, Claude API, Groq API, Tavily, Prompt Engineering, Structured Outputs (Pydantic)

Programming & Backend: Python, FastAPI, Streamlit, REST APIs, JSON

Data & Analysis: Pandas, NumPy, Matplotlib, Seaborn, Jupyter Notebook, EDA

Testing & Tools: pytest, unittest, Git, GitHub, MySQL, PostgreSQL, Linux, VS Code

ACHIEVEMENTS

- Certification: IBM RAG and Agentic AI Professional Certificate - IBM (Coursera).
- Deployed a live multi-agent Streamlit application (AI Mock Interview Coach) to Streamlit Community Cloud; built for the Razorpay Buildathon revenue-recovery track.
- Completed certifications in Python, SQL, and Machine Learning; active participant in AI/cloud workshops and hackathons.

EXTRA CURRICULARS

- Founded "BookVerse" reading club and coordinated alumni relations events; Video Editing Head, "Atrangi" Art & Craft Club.